

Personalization and recommendation workflow

This guide describes the implemented recommendation flow from segment configuration to the API response. It is intended for operators configuring a segment and engineers implementing, operating, or troubleshooting the workflow.

The production path runs through the `ai_agents_runners` Dagster master job. It is not the separate compatibility `personalization_job`, which remains a scaffold.

Scope and current behavior

The current workflow:

- selects active, tenant-scoped workflow steps in execution order from `cdp_agent_workflow`;
- supports API-triggered runs and cron-triggered runs;
- dispatches active `ranking_recommendation` agents for eligible segment members;
- ranks explicitly selected, tenant-owned `cdp_content_items` using tags, semantic similarity, or a hybrid score;
- persists run status and ranked items per profile; and
- serves the latest successful persisted results through the Customer 360 API.

Other model types are logged and skipped; this workflow does not dispatch them.

Personalization flow: end-to-end execution


At a business level, the configured workflow selects eligible profiles in a segment, ranks the chosen content for each profile, and makes successful recommendations available to the API. At runtime, the Dagster job coordinates the work asynchronously.

The orchestration is implemented in `dagster_defs.py` and `runner.py`. The shared model-type dispatcher is in `pipelines.py`; ranking logic is in `ranking_recommendation.py`.

1. **Trigger the run.** An API request supplies a tenant and optionally one segment. The cron sensor submits a run that checks eligible workflow schedules. Dagster assigns a run ID; recommendation records use this same ID.
2. **Build the execution plan.** The runner joins active `cdp_agent_workflow` steps to active registered `cdp_ai_agents`, applies the requested tenant and segment scope, and orders steps by `execution_order`. Cron runs use the workflow schedule override when present, otherwise the registered agent's schedule. This path executes only `ranking_recommendation` steps; other model types are logged and skipped.
3. **Select the audience.** For each tenant/segment pair, the runner loads active Customer 360 master profiles (the unified customer profiles) from `cdp_master_profiles` that belong to the active `cdp_segments` row. A profile belongs to the segment when its `segmentation_tags` contains the segment's `segment_tag`. The tenant ID scopes profile selection, candidate reads, and all writes; tenant ownership is rooted in `sys_tenant`.

4. **Prepare ranking inputs.** The `tags` strategy does not require embeddings. For `semantic` and `hybrid`, the runner checks that each selected candidate belongs to the tenant and is active. It generates or refreshes a canonical content vector when the source text, model, contract version, or vector dimensions have changed. It also embeds a profile query built from the profile's domain, segment tags, and optional `semantic_query`; it does not include personally identifiable information (PII). The configured embedding provider must produce 384- or 768-dimensional vectors.
5. **Rank content for each profile.** The runner passes the tenant, segment, agent, configuration, selected candidate IDs, profile domain and tags, and optional query vector to `execute_agent_pipeline`. The pipeline validates the payload and dispatches by `model_type` to `RankingRecommendationPipeline`. The handler rechecks tenant ownership, active status, selected IDs, and domain (`all` or the profile's domain). It scores candidates using tag overlap, semantic similarity, or a weighted combination; applies `minimum_score`; orders results consistently; and returns no more than `top_k` items.
6. **Persist the outcome.** The runner creates or resets a `RUNNING` row in `cdp_profile_recommendation_runs`, then upserts each profile's ranked items into `cdp_profile_recommendations`. On success, it marks the run `SUCCEEDED` and invalidates the tenant's recommendation cache. On failure, it marks the run `FAILED`, removes partial rows for that run, and propagates the error.
7. **Serve results.** `GET /api/v1/content-items/recommended` reads stored results; the request does not run ranking. The API checks Redis first. On a cache miss, it queries successful runs, merges eligible segment results, deduplicates content items, and returns the response. A successful recommendation run or relevant content mutation invalidates the tenant cache.


Runtime flow



The following sections describe workflow configuration, ranking behavior, persistence, and the recommendation read API. Each table name links to its `CREATE TABLE` definition in [database-schema.sql](#).

Manual test walkthrough

Use this walkthrough to verify the customer-facing flow and the backend execution path. The segment detail view and workflow editor are implemented in [segment-details.html](#) and [ai-agent-workflow.js](#).



Test procedure

1. **Confirm prerequisites.** Use an active segment with active profiles in the authenticated tenant. The segment's `segment_tag` must appear in each test profile's `segmentation_tags`. Confirm that the agent catalog has an active agent with `model_type = ranking_recommendation` and that the tenant has active content items.
2. **Open the workflow editor.** In the Segments page, open the target segment's details, select **Agent Workflow**, and click **Add agent**. Choose the active ranking agent, keep the step enabled, and set its execution order.

3. **Set ranking options and candidates.** For an initial smoke test, use `{"strategy": "tags", "top_k": 8}`; this does not require an embedding provider. Select at least one active content item or product. Each candidate must belong to the tenant and have domain `all` or the test profile's domain. Semantic and hybrid tests also require a supported embedding provider configured to generate 384- or 768-dimensional vectors.
4. **Save once to configure and run.** Click **Save workflow**. The UI sends `PUT /api/v1/segments/{segment_id}/workflow` with the complete ordered `steps` list. This is a full replacement of the segment workflow, not a partial update. The API validates and saves the configuration, then submits a Dagster run. The response includes `X-Dagster-Run-Id`; the UI displays **Saved · run queued**. Do not also call the run endpoint for the same test unless a second run is intended.
5. **Rerun without editing (optional).** To execute a saved workflow again, call `POST /api/v1/segments/{segment_id}/workflow/run` with the authenticated tenant context and a body such as `{"event": {"event_name": "personalization.manual_test"}}`. The API returns `202 Accepted` with `{"run_id": "...", "status": "submitted"}`. A manual API run is not held until the cron schedule; schedule matching applies to cron runs.
6. **Trace execution.** The API submits the run asynchronously to Dagster through `agent_workflow_api.py` and `dagster_client.py`. It does not call the Python ranking handler in the request process. Dagster starts `ai_agents_master_job` in the backend `ai_agents_runners` code location. The runner loads eligible profiles, prepares embeddings when required, and calls `execute_agent_pipeline` for each profile. The pipeline registry dispatches `ranking_recommendation` to `RankingRecommendationPipeline`, which applies the tenant, candidate, status, and domain filters before ranking.
7. **Verify the result.** Use `run_id` to find the run in the Dagster UI. `submitted` confirms that Dagster accepted the run; it does not mean the run has completed. After the run succeeds, call `GET /api/v1/content-items/recommended?master_profile_id=<profile_uuid>&segment_id=<segment_uuid>&limit=8` for an active profile in the segment. The response contains persisted recommendations for eligible items, or an empty list when none are eligible.

If saving returns `503` because Dagster could not accept the run, the workflow may already be saved; retry with `POST /api/v1/segments/{segment_id}/workflow/run`. If ranking fails, inspect the Dagster run and confirm the profile is an active segment member, candidates are active and tenant-owned, their domains match, and the selected strategy has the required embedding configuration.

Automated unit tests

Run the ranking handler tests from `customer360-backend`. The command measures statement and branch coverage for the handler and fails unless both reach 100%:

```
PYTHONPATH=../../customer360-dao/src pytest -q \  
ai_agents_runners/tests/test_ranking_recommendation.py \  
--cov=ai_agents_runners.agent_pipeline.agent_types.ranking_recommendation \  
--cov-branch --cov-report=term-missing --cov-fail-under=100
```

The test cases are in `test_ranking_recommendation.py`.

Agent registry: [cdp_ai_agents](#)

Agent definitions are global and use `agent_code` as their stable primary key. For workflow selection, the runner uses the agent code, `model_type`, `status`, and `schedule_definition`. Only agents with `status = 'ACTIVE'` are selected.

The registry also stores `model_name`, `input_features`, `hyperparameters`, `required_variables`, `system_instructions`, and prompt metadata. These fields are metadata for consumers that need them; the current ranking path does not pass them to its handler. In particular, `model_name` does not prove that a trained artifact is available. Runtime dispatch uses the implementation registered for the supported `model_type`.

The database supports `ranking_recommendation` as the model type used for content/product ranking. The seeded recommendation agent is `product_recommendation`; its current lifecycle status must still be checked at runtime rather than assumed from seed data.

Segment workflow: [cdp_agent_workflow](#)

Workflow rows are tenant-owned and attach one registered agent from [cdp_ai_agents](#) to one tenant-owned [cdp_segments](#) row under [sys_tenant](#).

Field	Purpose
<code>execution_order</code>	Positive, unique position within the tenant/segment workflow; lower values run first.
<code>is_active</code>	Controls whether the workflow step is eligible to run.
<code>schedule_definition</code>	Optional five-field cron override; if unset, the agent's schedule is used for cron runs.
<code>configuration</code>	JSON object containing strategy-specific settings.
<code>candidate_content_item_ids</code>	Selected content UUIDs for ranking; the API accepts up to 500 unique IDs per step.

The runner processes active steps in ascending `execution_order` and ignores inactive agents. API-triggered runs are immediate; schedule matching is applied to cron runs.

Candidate selection and ranking

Candidate catalog and tenant checks

Candidate IDs refer to [cdp_content_items](#). Each row belongs to a tenant and stores its item type (`news`, `video`, `product`, or `article`), display and call-to-action fields, segment tags, status, and publication metadata.

The workflow editor offers candidate selection only for `ranking_recommendation` agents. The API validates that each selected ID exists in the authenticated tenant; the database also enforces tenant ownership and distinct IDs. An active ranking step must contain at least one candidate. The runner evaluates only the selected IDs; it does not search the full catalog. At runtime, candidates must also be active and have domain `all` or the profile's domain.

Profile inputs and configuration

For each profile, the handler requires a non-blank `input_data.domain` and accepts `input_data.segmentation_tags` as a list of non-blank strings (default `[]`). Semantic and hybrid strategies also receive a profile embedding and an embedding model key from the runner. Profile context is built from the domain, segment tags, and optional `semantic_query`; it does not contain PII.

Setting	Default	Validation and behavior
<code>strategy</code>	<code>tags</code>	One of <code>tags</code> , <code>semantic</code> , or <code>hybrid</code> .
<code>top_k</code> (or <code>limit</code>)	<code>8</code>	Integer from 1 to 100; maximum result count.
<code>minimum_score</code>	<code>0</code>	Number from 0 to 1; lower-scoring candidates are excluded.
<code>semantic_weight</code>	<code>0.7</code>	Number from 0 to 1; used by <code>hybrid</code> .
<code>tag_weight</code>	<code>0.3</code>	Number from 0 to 1; used by <code>hybrid</code> . Hybrid weights must sum to more than 0.
<code>semantic_query</code>	<code>empty</code>	Optional text appended to the profile query; maximum 2,000 characters.

For `semantic`, the handler uses semantic weight 1 and tag weight 0. For `tags`, it uses tag weight 1 and semantic weight 0. The configured hybrid weights are normalized when the two scores are combined.

Scoring and result ordering

- **Tags:** the number of matched profile tags divided by the number of profile tags; the score is 0 when the profile has no tags.
- **Semantic:** $1 - \text{cosine distance} / 2$, clamped to the range 0–1.
- **Hybrid:** the weighted average of the tag and semantic scores. If the profile has no tags, the handler uses the semantic score when its weight is positive.

Candidates below `minimum_score` are excluded. Results are sorted by score descending, publication time descending (null values last), and content ID as the deterministic tie-breaker. Each result includes a one-based rank, the strategy scores, matched tags, a reason code, and the content display fields. Scores are strategy-specific ranking values, not calibrated probabilities.

Embedding lifecycle

Semantic and hybrid ranking require the configured `DOCS_EMBEDDING_PROVIDER` to be `openai` or `gemini`. The runner supports 384- and 768-dimensional vectors. For selected content, it generates or refreshes a vector when the canonical text, provider/model, contract version, or dimension is stale. Migration `008_content_embedding_contract.sql` preserves legacy 384- and 768-dimensional vectors and clears other sizes.

Products are ranked through their linked content items in the current workflow. Product-table vector columns are reserved for possible future direct product ranking.

Persistence and failure handling

The runner records run status in `cdp_profile_recommendation_runs` and per-profile ranked items in `cdp_profile_recommendations`. The read API serves these persisted records; it does not run ranking during the request.

If ranking produces no eligible results, the step fails rather than relaxing its filters or inventing a fallback. Failure diagnostics distinguish tenant-missing candidates, inactive content, profile-domain mismatch, unavailable embeddings, and scores below `minimum_score`. The error includes counts at each filter stage and the tenant, segment, agent, and profile identifiers. The detailed diagnostic query runs only on the failure path and uses the same tenant and selected candidate IDs.

Recommendation read API

The Customer 360 API exposes the persisted recommendations through:

```
GET /api/v1/content-items/recommended?master_profile_id=<uuid>&limit=8
```

`master_profile_id` is required. `limit` defaults to 8 and must be from 1 through 50. The API derives tenant scope from the authenticated request; a caller cannot request another tenant's data. Optional `segment_id` and `item_type` parameters restrict the result set.

For each profile, the API reads the latest successful run for each active segment the profile belongs to. It then merges results across those segments and keeps only the strongest recommendation for duplicate content items. Returned items must still match the profile domain and refer to active workflow, agent, and content records.

A missing or inactive profile returns `404`. A valid profile with no eligible persisted recommendations receives an empty list. Each result includes the content display fields and `matched_tags`, `segment_id`, `agent_code`, `rank`, `score`, `semantic_score`, `tag_score`, `strategy`, `reason`, and `generated_at`.

The API checks Redis before querying persisted results. Cache keys are scoped to the tenant, profile, and request filters. Content mutations and successful recommendation runs invalidate the tenant's recommendation cache; other changes become visible within `CACHE_TTL_SECONDS`. Set `CACHE_ENABLED=false` to bypass the cache while debugging.

The route is implemented in [content_api.py](#), and its tenant-scoped query is in [content_repository.py](#).

Current and future profile inputs

The current ranking contract reads only:

- the profile's domain and segmentation tags from [cdp_master_profiles](#);
- segment membership from [cdp_segments](#); and
- explicitly selected content candidates from [cdp_content_items](#).

It does not currently use persona-match records from [cdp_customer_personas](#) or shared persona definitions from [cdp_persona_archetypes](#). Additional analytics or persona fields stored on a master profile are also outside the current ranking input contract.

Any future handler that adds these inputs must define tenant-scoped reads and apply relevant consent, suppression, inventory, campaign-eligibility, and domain constraints before returning a recommendation.

Data ownership and engineering guardrails

The ranking handler validates its profile inputs and configuration, evaluates only the step's selected candidates within the step tenant, and returns ranked items. The runner—not the ranking handler—persists per-profile results under the Dagster run ID and updates run status.

The workflow must not:

- read or write data across tenant boundaries;
- rewrite identity links or source events;
- treat an unavailable handler or model as a successful default;
- rank candidates outside the configured, tenant-eligible set;
- send campaigns or notifications directly; or
- overwrite another handler's output without an explicit ownership contract.

References

- [AI runner orchestration](#)
- [Workflow schema](#)
- [Workflow migration](#)
- [Agent registry schema](#)
- [Content item model](#)